

# Introduction

# Team Members

**Lee Napthine**

Computer Science

**Parker DeBruyne**

Computer Science & Health  
Information Science

**Wesley Ducharme**

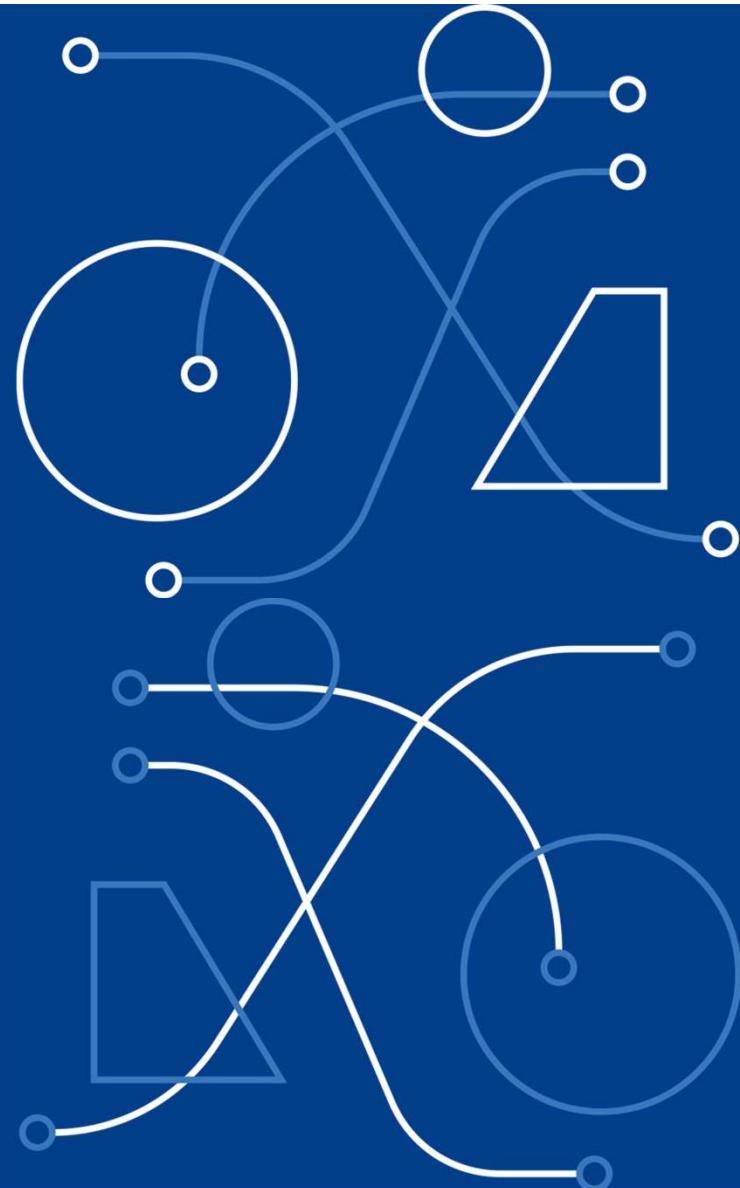
Computer Science &  
Statistics

**Shaban Shala**

Software Engineering

# Predicting Patient Visit Duration

via Synthetic Electronic Health Records



# Hospitals Struggle To Manage Staffing

## Bed Hours Occupied (BHO)

- **BHO = Discharge Time - Arrival Time**
- The total time a patient occupies a bed
- Summed over a period for all patients

## BHO Usage

- Analyze Capacity Management
- Operational Efficiency (bottlenecks)
- Resource Allocation
- Staffing & Finances
- Quality of Care
- Policy & Planning

## Without BHO Forecasts

- Inaccurate staffing
- Nurse overwork
- Resource strain
- Longer patient wait times
- Decreased Quality of Care



# Data Source: Synthea

## Synthetic Health Data

- Open-source software tool that generates synthetic, realistic electronic health records (EHRs)
- Simulates patient lifecycles and healthcare interactions

## Research & Development

- Allows researchers and developers to test and refine health IT applications
- Explore data science concepts
- Conduct simulations

## Privacy

- Entirely synthetic
- Anonymous
- i.e, No concerns with Data Privacy regulations

<https://synthetichealth.github.io/synthea/>



# Pivot: Patient Visit Duration

$m = 53,347$

Now:

- 53,347 Patient Encounters
- Enriched with individual Patient health data:
  - Age, Medications, Immunizations
- Enriched with Calendar data:
  - Day of the Week, Month, Year, Holidays
- Focused more on method and implementation



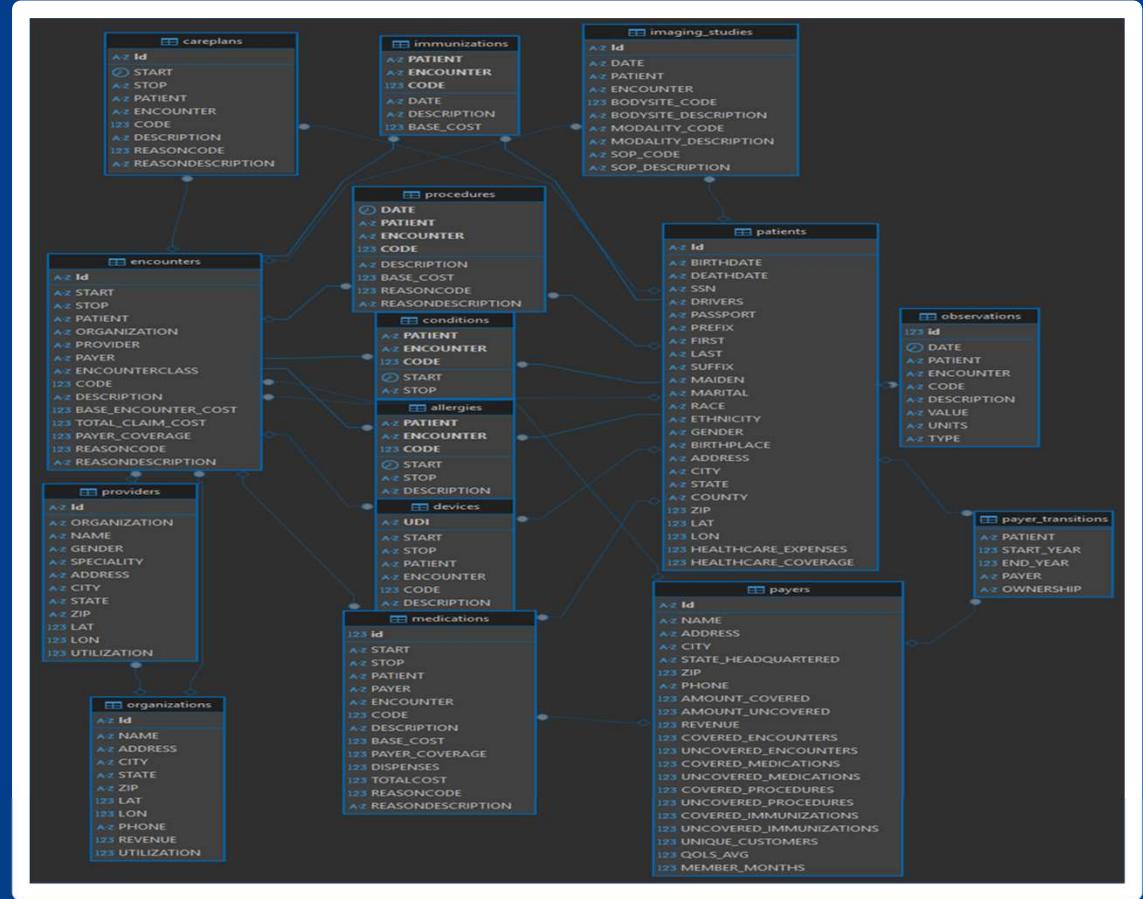
# Methods



# Setting up the Data

First steps:

- Made a PostgreSQL database to store the datasets for future querying.
- Hosted it on Railway for shared access.
- Queried the database to engineer our feature set and target variable. (Also engineered some features in our python scripts)
- Limited our data from 2010-2020.



# Features and Target

Initial feature set included:

- Description: Reason for the medical encounter
- Reason Description: Extra details about the encounter (if any)
- Gender
- Patient Age
- Organization Name
- And more ...
- Initial Target: Duration of Encounter (hours)

# Initial Models

Linear regression and SGD models with degree 2 polynomial features.  
Neural network with 1 hidden layer and 64 neurons. (RELU activation).

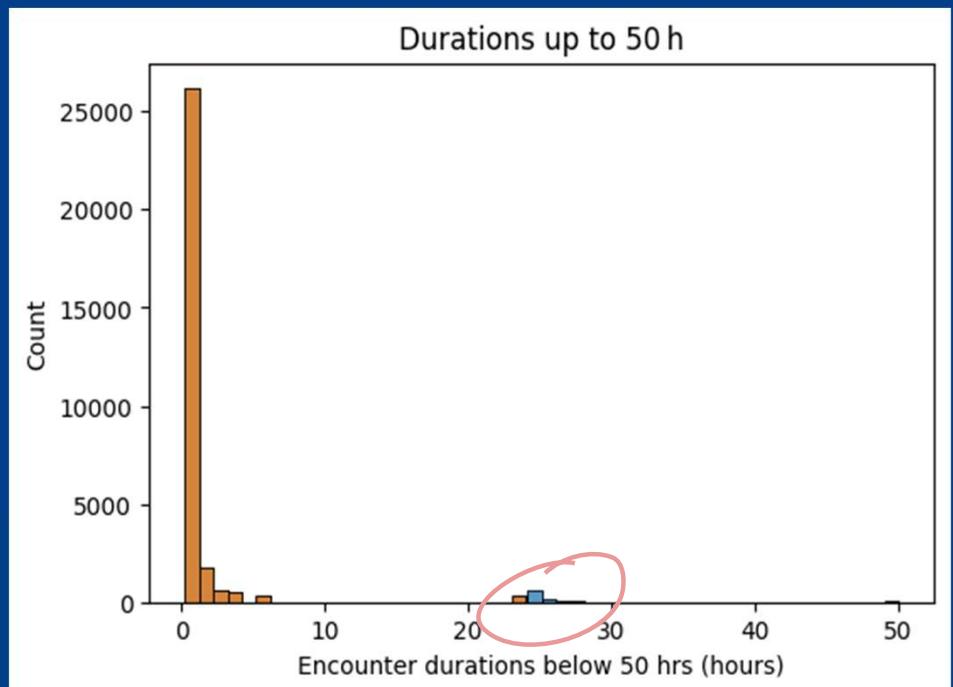
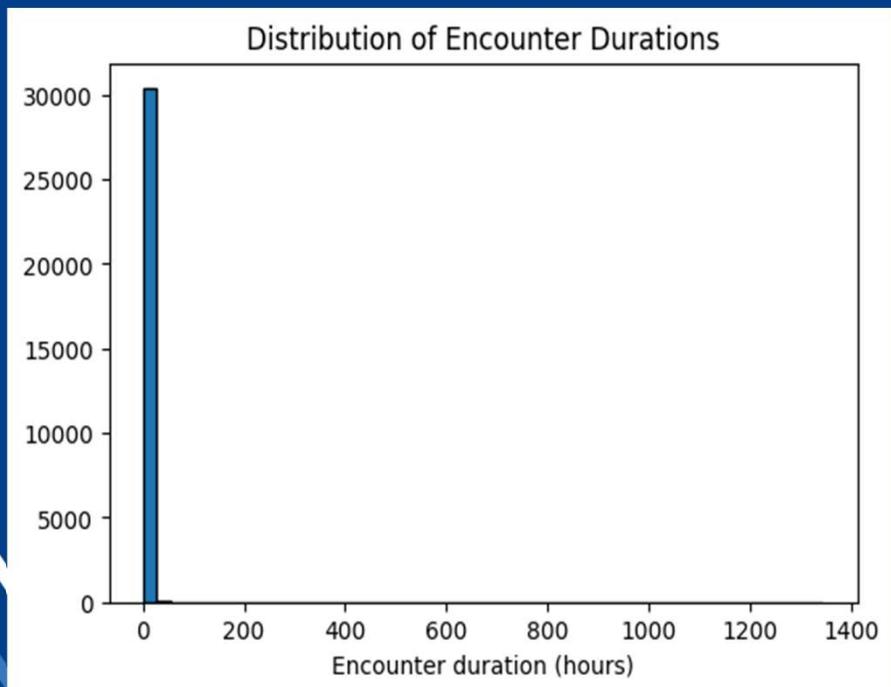
Both resulted in:

- very low  $R^2$ (approx 0.00001) metric
- very high RMSE(approx 10000+)
- log / sqrt transformations lead to no improvement

What could be causing problems?

Extreme outliers!

# Target Distribution



# New Approach

Swapped to a classification approach.

New target variable: `encounter_length_class`

Split it into 6 categories of duration length in hours:

- Quick (duration  $\leq$  0.25)
- Short (0.25 < duration  $\leq$  3)
- Moderate (3 < duration  $\leq$  8)
- Long (8 < duration  $\leq$  24)
- Very Long (24 < duration  $\leq$  50)
- Extreme (duration > 50)

Allowed us to both deal with the outlier problem while still giving accurate representation of the data.

# New Models

Now a classification problem!

New models that we will be comparing:

- Multinomial Logistic Regression model
- 3 Layer Neural Network with the new target





# Workflow and Debugging

Common workflow in both models:

- 80% training, 20% testing
- 3-Fold cross validation on the training set to fine tune hyper parameters.
- One last round of training and validation testing to view learning curves for possible overfitting and underfitting.
- After tuning do one final training and testing run on all training data and the unseen test data.

Debugging approach:

- Baseline was 70% accuracy and minimal overfitting / underfitting
- Check for Bias and Variance with learning curves
- Make changes appropriately to reduce overfitting / underfitting

# The Logistic Model

- Used Scikit-learn's built in Logistic Regression
- Set the class weight parameter to balance the weights
- Noticed highest accuracy with degree 2 polynomial features
- Initial learning curves showed high variance with no convergence
  - Introduced feature reduction to reduce this variance

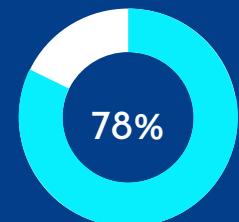
Early Implementation Learning Curve



# Logistic Regression Results

Test set results:

	precision	recall	f1-score
Quick ( $\leq 0.25h$ )	0.79	0.83	0.81
Short ( $0.25-3h$ )	0.82	0.70	0.75
Moderate ( $3-8h$ )	0.47	0.99	0.63
Long ( $8-24h$ )	0.85	0.89	0.87
Very Long ( $24-50h$ )	0.71	0.98	0.82
Extreme ( $>50h$ )	0.65	0.96	0.78
accuracy			
macro avg	0.71	0.89	0.78
weighted avg	0.79	0.78	0.78



# LR Learning Curve



# The Neural Network

- Used PyTorch to build the NN
- Balanced the class weights manually
- Early learning curves showed extreme variance as before
- Tried multiple layers and neurons assignments

Final setup: 3 hidden layers

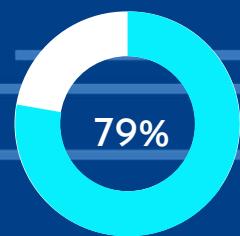
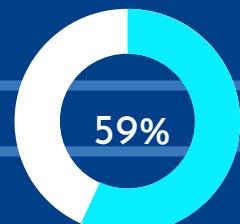
- First layer 128 Neurons
- Second layer 64 Neurons
- Third layer 32 Neurons
- RELU activation function
- Used dropout regularization ( $p\_drop=0.35$ )

(This being the final setup after having tested and checked multiple other layer and Neurons distributions gave the best results)

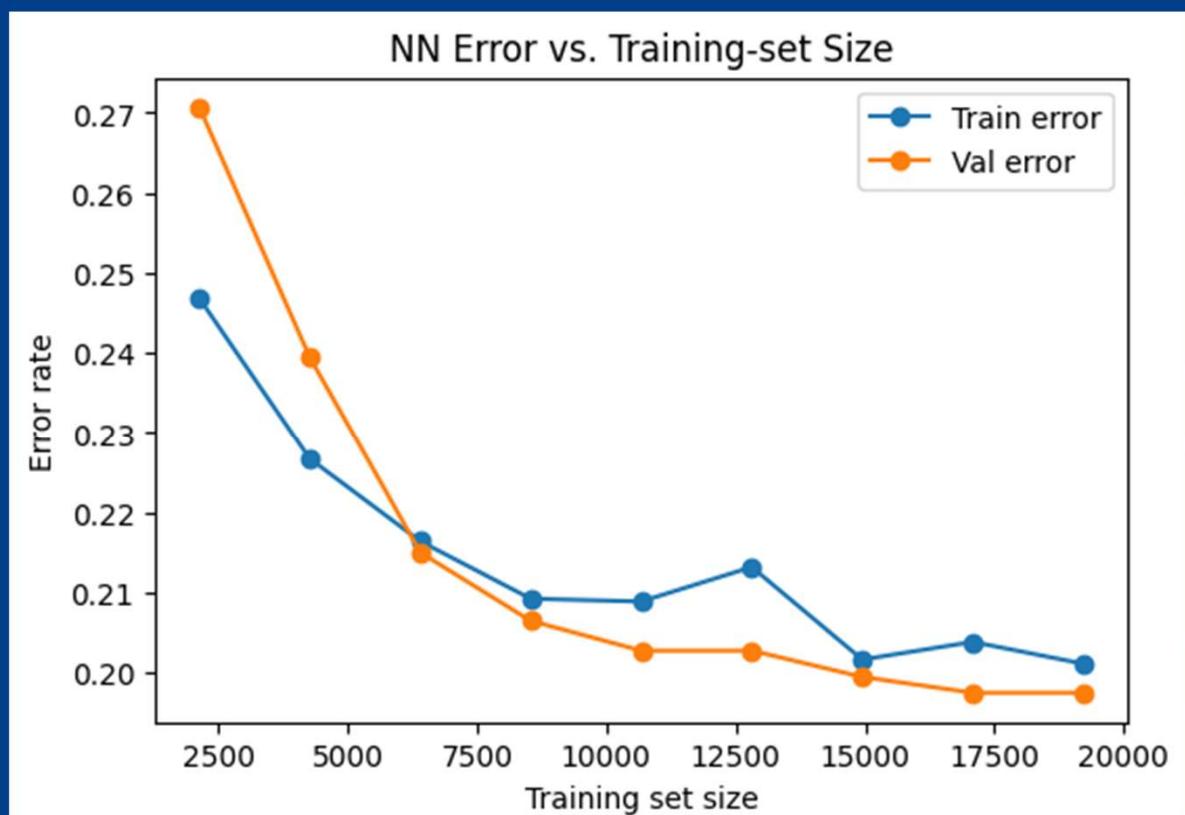


# Neural Network Results

	precision	recall	f1-score
Quick ( $\leq 0.25h$ )	0.73	0.96	0.83
Short ( $0.25-3h]$	0.96	0.59	0.73
Moderate ( $3-8h]$	0.50	0.98	0.66
Long ( $8-24h]$	0.81	0.96	0.88
Very Long ( $24-50h]$	0.88	0.98	0.93
Extreme ( $>50h)$	0.48	1.00	0.65
accuracy			
macro avg	0.73	0.91	0.78
weighted avg	0.83	0.79	0.78



# NN Learning Curve



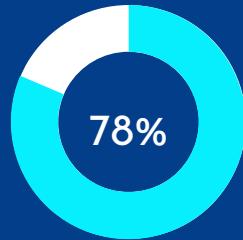
# Comparisons



# Results Comparisons

## Logistic Regression

Test set results:			
	precision	recall	f1-score
Quick ( $\leq 0.25h$ )	0.79	0.83	0.81
Short (0.25-3h]	0.82	0.70	0.75
Moderate (3-8h]	0.47	0.99	0.63
Long (8-24h]	0.85	0.89	0.87
Very Long (24-50h]	0.71	0.98	0.82
Extreme ( $>50h$ )	0.65	0.96	0.78
accuracy			0.78
macro avg	0.71	0.89	0.78
weighted avg	0.79	0.78	0.78



Logistic Regression Accuracy

	precision	recall	f1-score
Quick ( $\leq 0.25h$ )	0.73	0.96	0.83
Short (0.25-3h]	0.96	0.59	0.73
Moderate (3-8h]	0.50	0.98	0.66
Long (8-24h]	0.81	0.96	0.88
Very Long (24-50h]	0.88	0.98	0.93
Extreme ( $>50h$ )	0.48	1.00	0.65
accuracy			0.79
macro avg	0.73	0.91	0.78
weighted avg	0.83	0.79	0.78

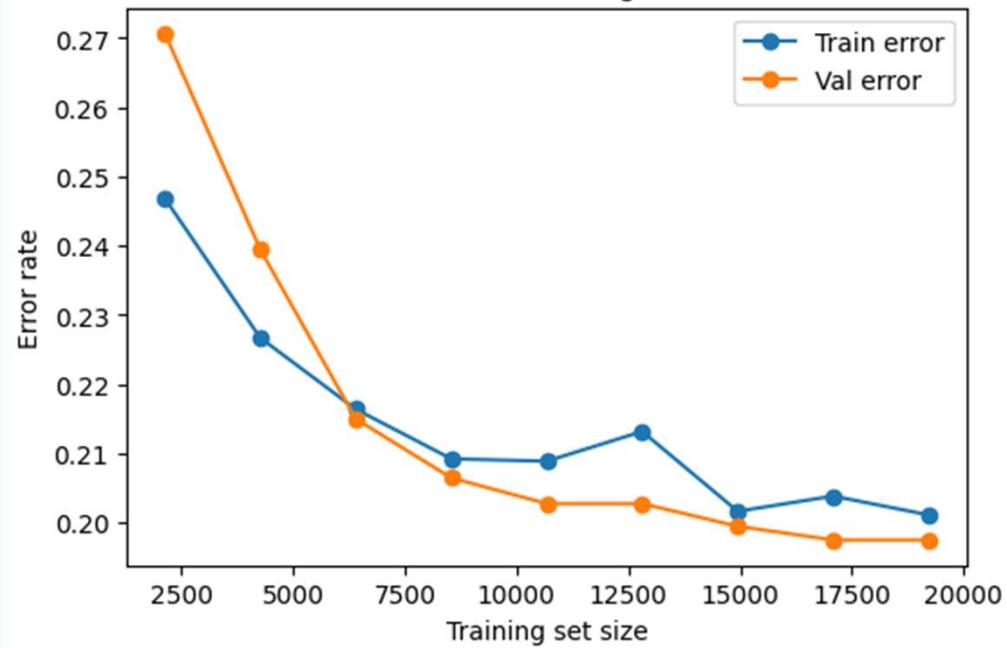


Neural Network Accuracy

# Visual Comparisons

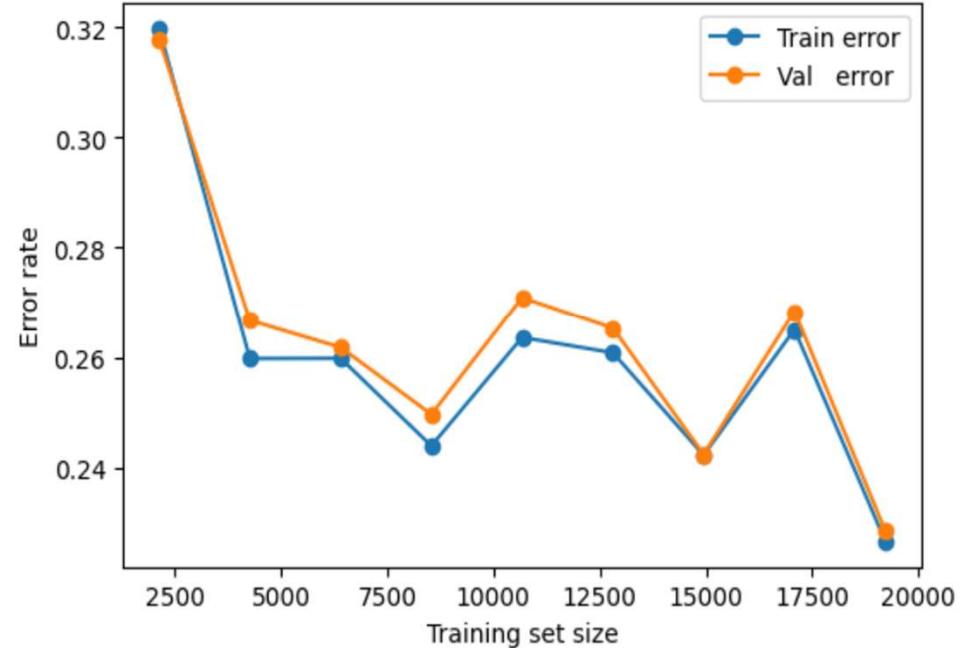
Neural Network

NN Error vs. Training-set Size



Logistic Regression

Logistic Regression - Error vs. Training-set Size



# Conclusions

## **Neural Networks captured patterns in our healthcare data with an impressive level of recall**

Recall was our prime metric, as missing on an outlier prediction could be costly.

## **Best performance came from cost bucket classification**

Reached **79% accuracy**, and our recall metrics were strong at predicting lengthy visits.

Model accuracy suffered when predicting short to moderate length visits with an average **54% accuracy**.



## **Data preprocessing and feature encoding are critical**

Highest weighted features: **visit type**, **visit description**, **medication costs**, and **age**.

These were predictably influential and reinforced our trust that the model was working.

## **Neural nets are powerful even with limited tuning**

With finer tuning our model could be capable of predicting encounter duration with enough accuracy for hospital planning.

# Limitations of Our Approach

## Limited time frame

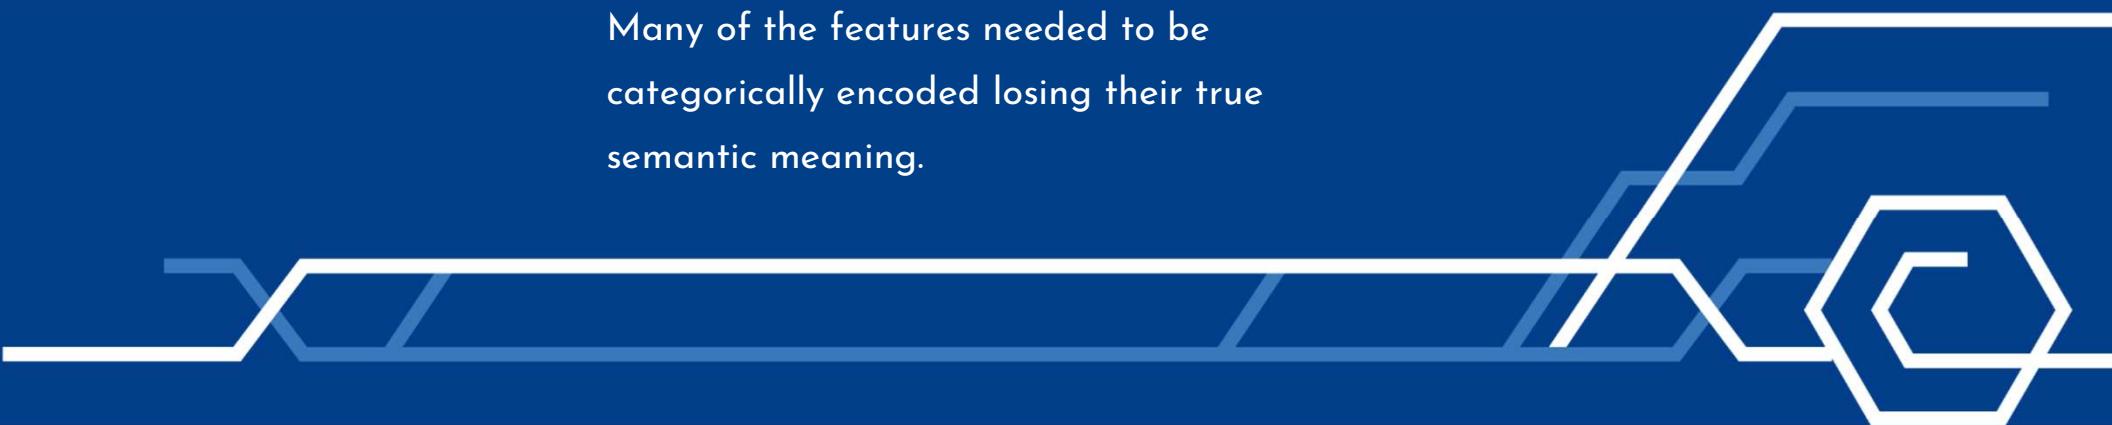
The dataset was very large and we could have spent more time engineering, testing, and iterating our features and model.

## Synthetic data

Very large synthetically generated dataset likely overstates the accuracy.

## Features grew too-large too-fast

Many of the features needed to be categorically encoded losing their true semantic meaning.





# Future Improvements

## Improve mid-range accuracy

The model struggled with mid-range visit durations. We plan to engineer additional features to address this gap and understand the causes

## Explore other types of NN's for time series cost forecasting

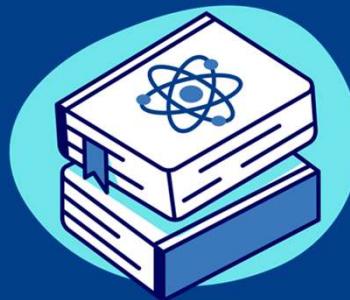
Certain studies had success using Recurrent Neural Networks for analysis on large medical datasets.

# Future Applications



## Hospital resource planning

Staff and bed  
scheduling



## Insurance analytics

Evaluating risk groups and  
costs



## Public health

Cost projections by  
region and Statistics  
Canada application



# Fairness and Bias

Feature engineering and model training may enforce biases that affect certain population over others

Ongoing evaluation is needed to prevent unfair outcomes

Thank you!

